


# Embedded Machine Learning Lab Challenge

By Valentin & Max

# Magnitude Pruning

pretrained model


- adapted for person class
- fine-tuned for 20 epochs  
-> 0.65 AP



# Magnitude Pruning


pretrained model

- adapted for person class
  - fine-tuned for 20 epochs
- > 0.65 AP



pruning  
~40% of  
params


after magnitude  
pruning  
-> 0.30 AP



# Magnitude Pruning


pretrained model

- adapted for person class
  - fine-tuned for 20 epochs
- > 0.65 AP




pruning  
~40% of  
params

after magnitude  
pruning  
-> 0.30 AP



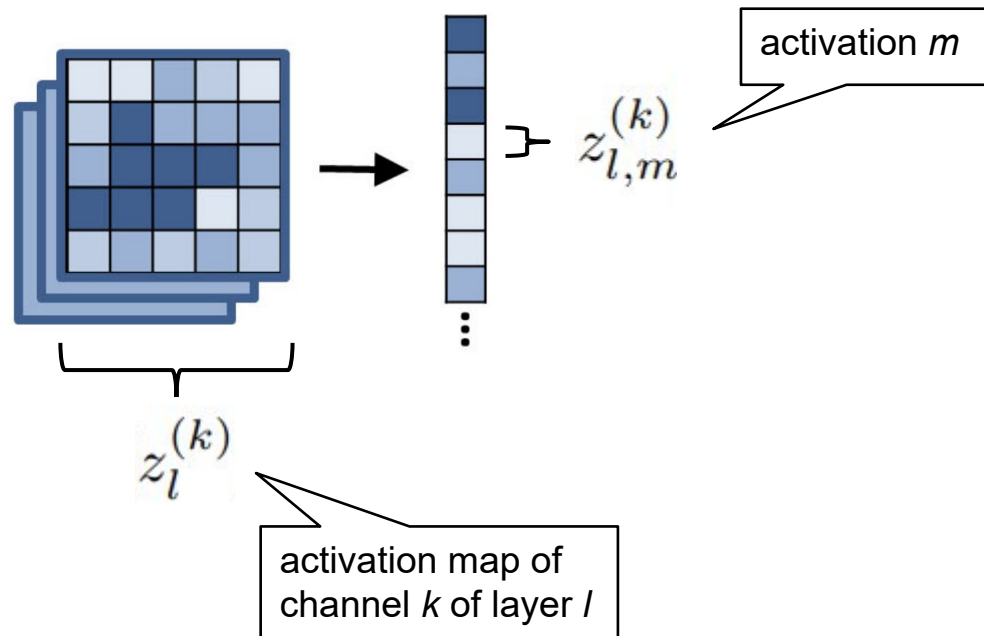
after Taylor  
pruning  
-> 0.51 AP



# Taylor Pruning <sup>[1]</sup>

## Channel pruning based on importance

- Defined by approximate change in loss caused by removing channel
- Activations & gradients gathered during forward passes

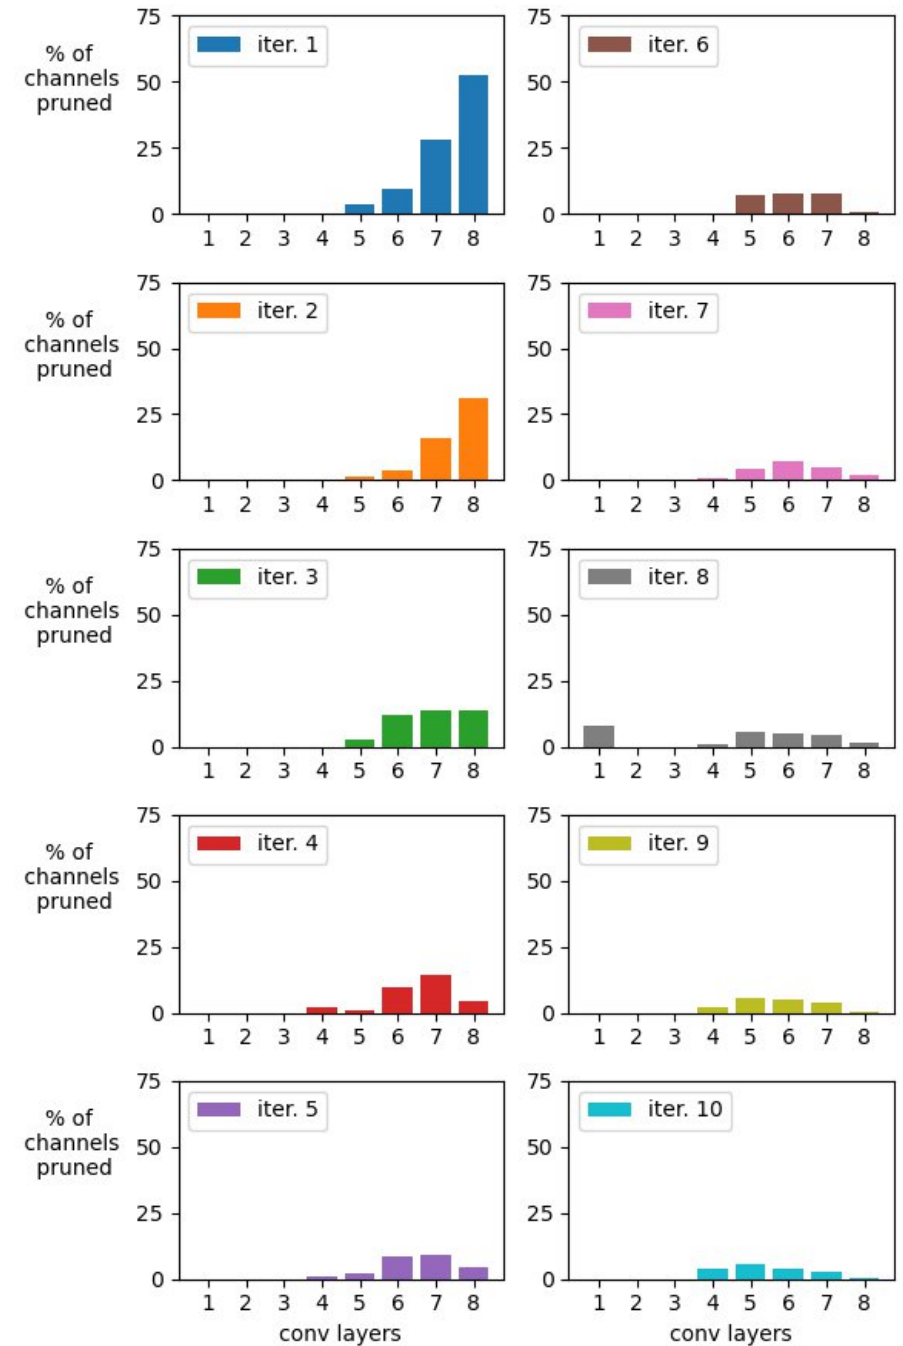


$$\underbrace{\Theta_{TE}(z_l^{(k)})}_{\text{importance of channel } k \text{ of layer } l} = \left| \frac{1}{M} \sum_m \frac{\delta C}{\delta z_{l,m}^{(k)}} z_{l,m}^{(k)} \right| \quad [1]$$


# Iterative Pruning

For 10 iterations:

1. Taylor prune  $k$  least important channels
  - $k$  proportional to num of params
2. Fine-tune for 10 epochs to regain performance
  - Very low learning rate
3. Evaluate AP
  - Compare APs over iterations



# Pruning Statistics



# Live camera flow

---

## Computation steps

1. capture image with camera (disabled JPEG encode for Jupyter)
2. run model on image
3. filter anchors, apply non-maximum suppression
4. draw bounding boxes and fps
5. compress image to JPEG (OpenCV should use libjpeg-turbo already)
6. show image to user (MJPEG stream)


**Jupyter notebook unreliable:** kernel crashes, unresponsive, ...

→ **run on a custom web server inside docker**




# Docker container

- provided image outdated (2021)
- use "l4t-ml" from dustynv (11/2023)  
<https://github.com/dusty-nv/jetson-containers>
- chose matching NVIDIA JetPack/L4T
- newer libraries, **faster inference**
- unchanged performance on other parts
- simple web server with Flask
- low resource usage
- has endpoint serving MJPEG stream
- UI to start/stop camera



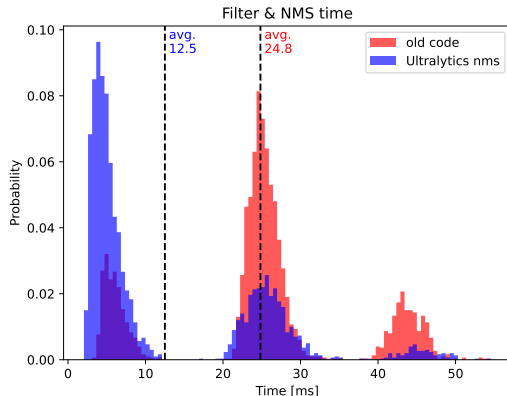
# ONNX runtime

- use for inference on GPU
- supposedly faster than PyTorch
- represents model computations as graph
- use optimization `ORT_ENABLE_ALL`  
→ fuses conv and batch-norm layers
- Quantization
  - int8 → not supported by GPU
  - float16 → no speed up



# Filter boxes and NMS

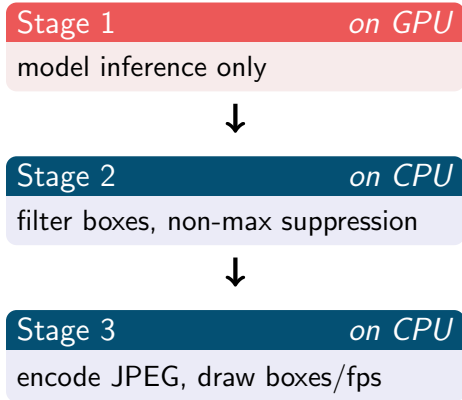
- original code: slow, pure Python impl.
- use `non_max_suppression()` from Ultralytics library
- licensed as AGPL v3
- everything combined in one function
- based on native `torchvision.ops.nms()`



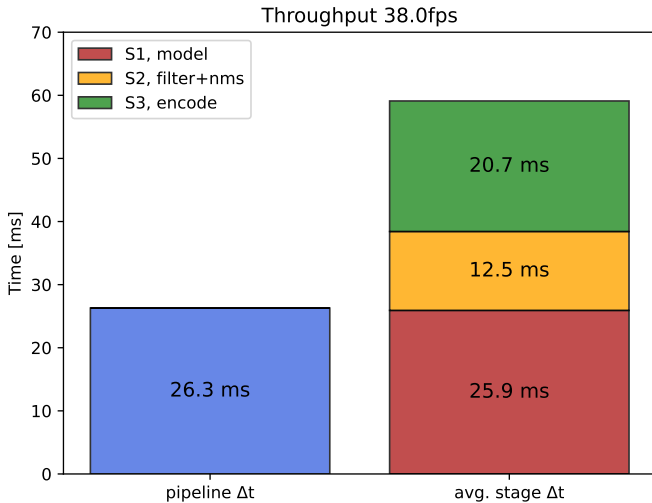
# Pipelining

---

- split work between GPU and CPU
- parallelize steps
- as `Python threading.Thread`  
→ GIL mostly no issue
- connected by queues, length limited
- **Input:** image from camera callback
- **Output:** byte stream for HTTP response



# Pipelining: results



# References

- [1] Pavlo Molchanov, Stephen Tyree, Tero Karras, Timo Aila, & Jan Kautz. (2017). Pruning Convolutional Neural Networks for Resource Efficient Inference.