

Live camera flow

Computation steps

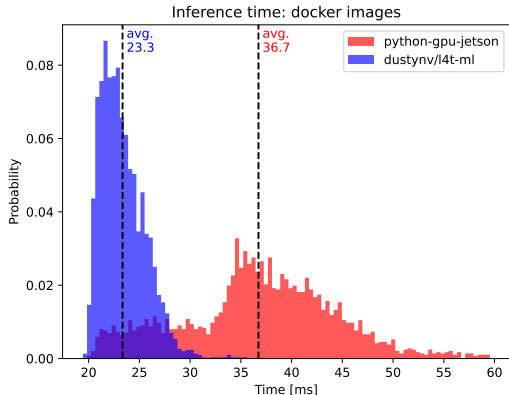
1. capture image with camera (disabled JPEG encode for Jupyter)
2. run model on image
3. filter anchors, apply non-maximum suppression
4. draw bounding boxes and fps
5. compress image to JPEG (OpenCV should use libjpeg-turbo already)
6. show image to user (MJPEG stream)

Jupyter notebook unreliable: kernel crashes, unresponsive, ...

→ **run on a custom web server inside docker**

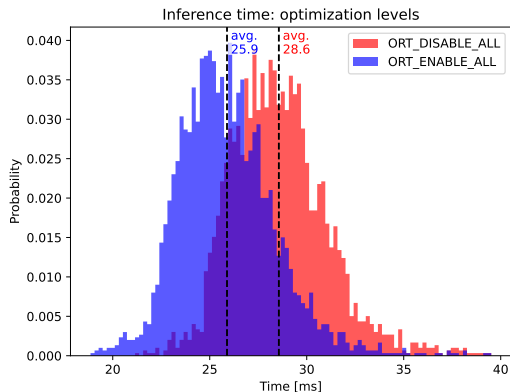
Docker container

- provided image outdated (2021)
- use "l4t-ml" from dustynv (11/2023)
<https://github.com/dusty-nv/jetson-containers>
- chose matching NVIDIA JetPack/L4T
- newer libraries, **faster inference**
- unchanged performance on other parts
- simple web server with Flask
- low resource usage
- has endpoint serving MJPEG stream
- UI to start/stop camera



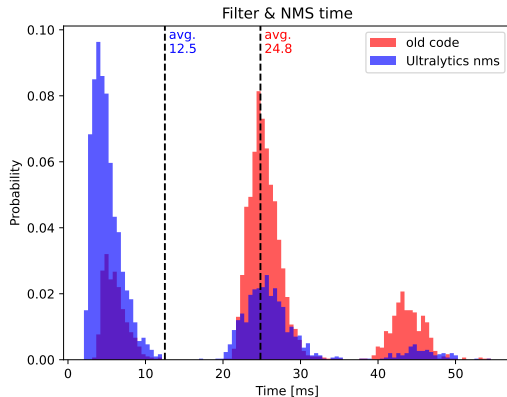
ONNX runtime

- use for inference on GPU
- supposedly faster than PyTorch
- represents model computations as graph
- use optimization `ORT_ENABLE_ALL`
→ fuses conv and batch-norm layers
- Quantization
 - int8 → not supported by GPU
 - float16 → no speed up



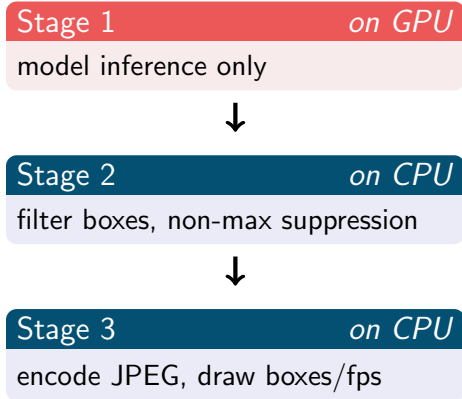
Filter boxes and NMS

- original code: slow, pure Python impl.
- use `non_max_suppression()` from Ultralytics library
- licensed as AGPL v3
- everything combined in one function
- based on native `torchvision.ops.nms()`



Pipelining

- split work between GPU and CPU
- parallelize steps
- as `Python threading.Thread`
→ GIL mostly no issue
- connected by queues, length limited
- **Input:** image from camera callback
- **Output:** byte stream for HTTP response



Pipelining: results

